



PRESENTS

THE AI READINESS FILES

Developer Reference Guide
Getting Your Website Ready for the Age of AI

17
FILES

7
CATEGORIES

DEPLOY
READY



Developer Overview

This reference covers all AI Readiness Files — structured assets deployed to a website's root that enable AI systems (LLMs, crawlers, RAG pipelines) to accurately index, understand, and represent the site's content. Each file targets a specific layer of the AI consumption stack: crawler access control, content context, structured entity data, retrieval indexing, and policy declaration.

Deploy Critical files first. All URLs referenced in any file must resolve to HTTP 200. Validate JSON files before deployment and test robots.txt changes against AI crawler user-agent strings. See deployment-checklist.md for the full phased rollout sequence.

Category	Files	Primary Consumers
Identity & Permissions	robots.txt, ai.txt	All AI crawlers; LLM training pipelines
Content Files	llms.txt, llms-full.txt	LLMs (ChatGPT, Claude, Gemini, Perplexity); RAG pipelines
Map & Navigation	ai-sitemap.xml, sitemap.md	AI crawlers; semantic search indexers
Intelligence Files	ai-entities.json, ai-intent.json, ai-schema.json + ai-plugin.json	Knowledge graphs; conversational AI; search engines
Research Files	rag-index.json, rag-index.jsonl	LlamaIndex, LangChain, Pinecone, Weaviate, OpenAI fine-tuning
Policy Files	ai-disclosure.txt, training-data-policy.txt	AI model developers; regulatory auditors
Operations Files	structured-data-guide.md, manifest.json, deployment-checklist.md	Dev team; DevOps; partner integrations

Deployment Priority

Priority	Files	Deploy When
Critical	robots.txt, ai.txt, llms.txt, llms-full.txt	Before any AI indexing begins
High	ai-sitemap.xml, ai-entities.json, ai-intent.json, ai-schema.json, rag-index.*	Sprint 1 — within first deployment
Medium	Policy files, sitemap.md, manifest.json, .well-known/ai-plugin.json	Sprint 2 — within 30 days
Internal	structured-data-guide.md, deployment-checklist.md	Internal use; optionally serve publicly

IDENTITY & PERMISSIONS

IDENTITY & PERMISSIONS

robots.txt [Critical]

/robots.txt

Format: text/plain **Update:** On structural change

Extend your existing robots.txt with per-agent User-agent blocks for each major AI crawler. Grant Allow access to all AI readiness files (llms.txt, ai-sitemap.xml, rag-index.*) while Disallowing transactional paths (/cart/, /checkout/, /account/). Add Sitemap: directive pointing to both standard and ai-sitemap.xml. Avoid CMS-specific Disallow rules unless confirmed.

CONSUMED BY

GPTBot, ClaudeBot, Google-Extended, PerplexityBot, CCBot, all crawlers

IDENTITY & PERMISSIONS

ai.txt [Critical]

/ai.txt

Format: text/plain **Update:** Quarterly

Root-level brand declaration for AI systems. Declare organization name, industry, canonical URL, primary product taxonomy, authoritative topic list, and training data policy (RAG-permitted vs. license-required for commercial model training). Functions as the AI-era equivalent of robots.txt. Keep factual, structured, and version-stamped.

CONSUMED BY

AI crawlers, LLM training pipelines, RAG indexers

CONTENT FILES

CONTENT FILES

llms.txt [Critical]

/llms.txt

Format: text/plain (Markdown) **Update:** Monthly

Follows the llmstxt.org spec: H1 org name, blockquote description, H2 category sections with markdown links, priority queries at the bottom. Keep under ~4,000 tokens to fit a single LLM context window. All URLs must resolve to confirmed 200 responses. Primary discovery file for AI assistants doing site lookups.

CONSUMED BY

LLMs: ChatGPT, Claude, Gemini, Perplexity; AI search engines

CONTENT FILES

llms-full.txt [Critical]

/llms-full.txt

Format: text/plain (Markdown) **Update:** Quarterly

Extended context file for RAG and knowledge-base ingestion. Includes per-category rich descriptions, FAQ/Q&A; pairs, and authoritative answers to top-intent queries. No token limit enforced — optimised for retrieval-augmented pipelines that chunk content. Mark time-sensitive data (pricing, inventory) with explicit staleness warnings.

CONSUMED BY

RAG pipelines, AI knowledge bases, Perplexity, deep-indexing LLMs

MAP & NAVIGATION

MAP & NAVIGATION

ai-sitemap.xml [High]

/ai-sitemap.xml

Format: application/xml **Update:** Monthly

XML sitemap extended with custom ai: namespace tags per URL: ai:content-type, ai:topic, and ai:summary. Priority values 1.0 (homepage) down to 0.5 (secondary pages). Reference from robots.txt Sitemap: directive and submit to Google Search Console alongside the standard sitemap. Validate against the sitemaps.org protocol before deployment.

CONSUMED BY

AI crawlers, semantic search indexers, RAG pipelines

MAP & NAVIGATION

sitemap.md [Medium]

/sitemap.md

Format: text/markdown **Update:** Monthly

Human and machine-readable site map in plain Markdown. Organised into H2 sections matching primary site categories, each with a markdown link list. No XML parsing required — ideal for LLMs processing the site via text context rather than structured crawl. Keep in sync with ai-sitemap.xml.

CONSUMED BY

LLMs, developer tooling, content teams

INTELLIGENCE FILES

INTELLIGENCE FILES

ai-entities.json [High]

/ai-entities.json

Format: application/json **Update:** Quarterly

JSON entity declaration file. Schema: array of objects with fields name, type (Product, Service, Brand, Category, etc.), description, url, and related_entities[]. Powers knowledge graph construction and entity disambiguation. AI systems use this to confidently classify offerings and resolve naming variants across the web.

CONSUMED BY

Knowledge graph builders, entity resolution engines, AI systems

INTELLIGENCE FILES

ai-intent.json [High]

/ai-intent.json

Format: application/json **Update:** Quarterly

Query-to-URL intent mapping. Schema: array of objects with query (natural language), intent_type (informational | navigational | transactional | comparison), and preferred_url. AI assistants use this to route users to the correct landing page. Cover 30–50+ intents. All preferred_url values must be confirmed live URLs.

CONSUMED BY

Conversational AI, AI search engines, AI assistants

INTELLIGENCE FILES

ai-schema.json [High]

/ai-schema.json

Format: application/json (JSON-LD) **Update:** Annually

Standalone JSON-LD file using Schema.org vocabulary. Primary @type: Organization + Store (or EducationalOrganization, SoftwareApplication per classification). Include sameAs social links, foundingDate, address, contactPoint, and SearchAction for site search. Mirrors the inline JSON-LD in the page head. Validate at validator.schema.org.

CONSUMED BY

Search engines, AI systems, entity databases, developers

INTELLIGENCE FILES

.well-known/ai-plugin.json [Medium]

/well-known/ai-plugin.json

Format: application/json **Update:** On contact/category change

OpenAI-compatible plugin manifest served from the .well-known directory. Ensure Content-Type: application/json and CORS headers allow cross-origin access. Required fields: schema_version, name_for_human, name_for_model, description_for_human, description_for_model, auth type, and contact_email. Create the .well-known/ directory if not already present.

CONSUMED BY

AI assistant platforms, plugin registries, developer platforms

RESEARCH FILES

RESEARCH FILES

rag-index.json [High]

/rag-index.json

Format: application/json **Update:** Monthly

JSON array of indexable records. Each object: { url, title, topics: string[], summary?, lastModified? }. One record per major page or logical content section. Pipeline engineers ingest this directly with no additional transformation. Auto-generate rag-index.jsonl from this file. All URLs must be live — dead links degrade retrieval quality.

CONSUMED BY

LlamaIndex, LangChain, Pinecone, Weaviate, custom RAG pipelines

RESEARCH FILES

rag-index.jsonl [High]

/rag-index.jsonl

Format: application/x-ndjson **Update:** Sync with rag-index.json

Newline-delimited JSON (NDJSON) version of rag-index.json — identical records, one per line. Required for OpenAI fine-tuning datasets and most vector database bulk-import tools. Generate via: python3 -c "import json; [print(json.dumps(r)) for r in json.load(open('rag-index.json'))]". Set Content-Type: application/x-ndjson on server.

CONSUMED BY

OpenAI fine-tuning, Pinecone, Weaviate, NDJSON-compatible pipelines

POLICY FILES

POLICY FILES

ai-disclosure.txt [Medium]

/ai-disclosure.txt

Format: text/plain **Update:** Annually

Plain-text AI usage disclosure. Declare which AI tools are used in production (chatbots, recommendations, content generation, etc.), data handling practices, human oversight mechanisms, and complaint/escalation contacts. AI systems read this to accurately report your AI posture. Reference from your privacy policy and Terms of Service.

CONSUMED BY

AI systems, regulatory audits, transparency-checking tools

POLICY FILES

training-data-policy.txt [Medium]

/training-data-policy.txt

Format: text/plain **Update:** Annually

Formal content licensing policy for AI training use cases. Distinguish permitted uses (RAG retrieval, search indexing) from restricted uses (commercial model training, dataset resale). Include staleness warnings for time-sensitive content. Complements robots.txt — robots.txt handles crawl access, this handles downstream use rights.

CONSUMED BY

AI model developers, data licensing teams, legal/compliance tools

OPERATIONS FILES

OPERATIONS FILES

structured-data-guide.md

/structured-data-guide.md

Format: text/markdown **Update:** On new page-type launch

Developer guide for Schema.org JSON-LD across all page types: Organization, Store, Product+Offer, ItemList, BreadcrumbList, LocalBusiness, Event, BlogPosting, WebSite+SearchAction. Includes copy-paste code blocks per template. Validation workflow: Google Rich Results Test and validator.schema.org. Serve publicly or keep internal.

CONSUMED BY

Internal dev team, SEO engineers, AI readiness auditors

OPERATIONS FILES

manifest.json [Medium]

/manifest.json

Format: application/json **Update:** On file set changes

Machine-readable inventory of the full AI readiness file set. Fields per entry: filename, deploy_path, url, content_type, purpose, audience[], update_frequency, priority. Totals section: file count, priority breakdown, RAG record count, intent count. Enables automated compliance checks and partner discovery of your implementation.

CONSUMED BY

Dev tooling, AI readiness auditors, partner integrations

OPERATIONS FILES

deployment-checklist.md

/deployment-checklist.md

Format: text/markdown **Update:** On file set changes

Phased deployment playbook: Phase 1 critical text files, Phase 2 sitemaps, Phase 3 JSON data files, Phase 4 policy files, Phase 5 .well-known/ setup, Phase 6 robots.txt update, Phase 7 JSON-LD implementation, Phase 8 Search Console submission. Includes HTTP 200 verification commands and maintenance schedule.

CONSUMED BY

Dev team, DevOps, project managers

Schema.org Type Reference

Select the correct Schema.org @type for ai-schema.json and inline JSON-LD. Multiple types can be combined using @graph.

Site Type	Primary @type	Additional Types
B2B SaaS / Software	SoftwareApplication	Organization
E-commerce / Retail	Store	Organization, Product, Offer
Education / University	EducationalOrganization	EducationalOccupationalProgram
Healthcare	MedicalOrganization	Organization, Physician
Professional Services	Organization	LocalBusiness, Service
Nonprofit / NGO	NGO	Organization, Event
Media / Publishing	Organization	WebSite, BlogPosting, NewsArticle
Restaurant / Hospitality	Restaurant	LocalBusiness, Menu

Maintenance Schedule

File(s)	Frequency	Trigger
llms.txt / llms-full.txt	Monthly / Quarterly	Product, service, or pricing changes
ai-intent.json	Quarterly	New high-intent queries in analytics
ai-entities.json	Quarterly / Launch	New product categories or services
ai-sitemap.xml / sitemap.md	Monthly	New or removed major pages
rag-index.json + .jsonl	Monthly	New content sections or URL changes
ai-schema.json	Annually	Company info, address, social link changes
training-data-policy / ai-disclosure	Annually	Legal review or AI tool usage changes
robots.txt / ai.txt	On change	New AI crawlers or site structure changes